

ALEJANDRO ORTIZ PERDOMO

AI Engineer | Machine Learning Engineer

📍 Bogotá, Colombia

☎ +57 319 666 3193

✉ alejortizp@gmail.com

linkedin.com/in/alejandro-ortiz-perdomo |

github.com/alejortizp |

🌐 alejortizp.github.io/curriculum_HV

PROFESSIONAL SUMMARY

MLOps Engineer specialized in taking ML models from experimentation to **reliable production**. Implemented **CI/CD** pipelines with **Databricks + MLflow** reducing deployment from 2 weeks to **2 days**. Designed **data/model drift** monitoring systems with automated retraining alerts. Proficient in **Docker**, **Azure AI**, **FastAPI** and **GitHub Actions** for scalable and reproducible ML infrastructure.

EDUCATION

Master in Software Engineering

[Universidad de los Andes](#)
2025 – In Progress

B.S. in Statistics

[Universidad del Valle](#)
2011 – 2016

Data Science Diploma (DS4A)

Correlation One, 2022

Deep Learning Diploma

Univ. Autónoma Occidente, 2023

LANGUAGES



Spanish

Native



English

Professional Proficiency
(B2)

KEY SKILLS

Effective Communication

Problem Solving

Analytical Thinking

Team Collaboration

Adaptability

Continuous Learning

Time Management

Technical Leadership

WORK EXPERIENCE

AI Engineer — AI Practice

Periferia IT Group

Jun 2025 – Present

- **PerXIA Suite Ecosystem (SRPA AI Agents):** Contribute to the technical design, conceptualization and documentation of Periferia's proprietary *Single-Responsibility Plug-and-play Agents* platform automating the end-to-end software development lifecycle. Key contributions to **PERXIA-IDEATOR** (developer onboarding via repository knowledge graphs, cutting ramp-up by ~70%), **PERXIA-HUS** (automated user-story generation with acceptance criteria), **PERXIA-DEV** (Cursor IDE contextual assistance), **PERXIA-UNIT** (multi-language unit testing with ≥85% coverage), **PERXIA-QA**, **PERXIA-DOCS** and **PERXIA-HELP** (assistant built on **DeepSeek R-1** deployed on **Azure AI Foundry**). Multi-cloud deployment (**Azure / AWS / GCP**) delivering **60–70%** savings across the development cycle.
- **CV-Job Matching System:** Pipeline with **Azure Document Intelligence + AI Search** processing 500+ CVs/day with **95%** extraction accuracy and vector semantic search reducing screening time by **70%**. Running in production at **10k+ searches/month** with <200ms latency.
- **Pre-sales & Solution Architecture:** Build technical estimates, team sizing and reference architectures for **GenAI, Agentic AI and RAG** solutions for enterprise clients (banking, insurance, retail, energy and telco). Lead technical discovery with stakeholders and position PerXIA Suite agents in commercial technical proposals.

Machine Learning Engineer

Global Hitss

Sep 2023 – Jun 2025

- **GenAI & RAG:** Designed and deployed conversational agent with RAG that reduced support tickets by **40%**, processing **10k+ queries/month** with 85% resolution without escalation. Stack: LangChain, Azure OpenAI, FAISS.
- **MLOps:** Implemented CI/CD pipeline in **Databricks + MLflow** reducing deployment time from 2 weeks to **2 days**. Production models with automatic drift monitoring and retraining.
- **Reinforcement Learning:** Productized *Next Best Offer* (NBO) system with A/B testing that increased conversion by **15%** and revenue by **\$10k per person**.
- **Churn Model:** Developed predictive model with **Lift of 3.0** and 82% recall in top decile, integrated to CRM for proactive retention. Impacted **50k+ customers**.

CERTIFICATIONS

AI-102: Microsoft Certified Azure AI Engineer Associate

Microsoft, Feb 2026

NLP Specialization (RAG/Transformers)

Coursera, Jun 2025

PyTorch Ultimate

Coursera, Mar 2025

MLOps and Databricks

Netec, Dec 2024

AWS Cloud Practitioner

Udemy, Feb 2022

TECHNICAL SKILLS

(years exp.)

GENAI & LLMS

LangChain 2+ (Agents, RAG)

Azure OpenAI 2+ (GPT-4, Embeddings)

Prompt Engineering 2+

Vector DBs 2+ (FAISS, Pinecone)

AZURE AI & CLOUD

Azure AI Search 1+

Document Intelligence 1+

Azure AI Foundry 1+

Databricks + MLflow 3+

GOOGLE CLOUD PLATFORM (GCP)

Vertex AI 1+ (Training, Endpoints, Pipelines)

Cloud Run 1+

Compute Engine 1+ (MIG zones)

Cloud Storage 1+ | Cloud SQL 1+

ML/DL FRAMEWORKS

PyTorch 2+ | TensorFlow 3+

Scikit-learn 4+ | XGBoost 4+ | LightGBM 4+

Hugging Face Transformers 2+

MLOPS & ENGINEERING

Python 7+ | SQL 7+ | PySpark 3+

Docker 2+ | CI/CD 2+

FastAPI 3+ | Streamlit 3+

Git 4+ | GitHub Actions 2+

Machine Learning Consultant

Aug – Sep 2023

Grupo CINTE

- **ML Observability:** Designed comprehensive monitoring methodology for *data/model drift* detection with automated retraining alerts, adopted as the data science team's standard.
- **Optimization:** Audited and optimized LightGBM and Random Forest algorithms, improving computational efficiency and reducing inference times.

Data Scientist

Mar 2023 – Jul 2023

NAE

- **Contactability Model:** Developed **LightGBM** model that increased effective contact rate by **60%** (from 25% to 40%), optimizing time windows and channels. Impacted campaigns with **100k+ contacts/month**.

Data Scientist II

May 2022 – Jan 2023

Xolit

- **Payment Propensity Model:** Developed predictive model for payment probability and portfolio agreement, maximizing collection through **automated management priorities** based on propensity scores.
- **Recommendation System:** Implemented optimal contact time-slot recommendation engine per account holder, increasing debt collection management effectiveness.
- **Customer Segmentation:** Designed clustering-based marketing strategy for account holders, enabling personalized communication per segment to optimize collection rates.
- **Monitoring:** Built **Power BI** dashboards for tracking collection KPIs and production model performance metrics.

OPEN SOURCE PROJECTS

Email Agent with LangChain

— github.com/alejortizp/email-agent

Intelligent agent that queries emails via IMAP, classifies with LLM, and generates contextual automatic responses. Reduces response time by **80%**.

Stack: LangChain, Azure OpenAI, IMAP, Python, Prompt Engineering

NL-to-SQL/Pandas Agent — github.com/alejortizp/nl-to-code

Agent that translates natural language to SQL queries and Pandas operations. Includes code validation, secure sandbox execution, and automatic results visualization.

Stack: LangChain, SQL, Pandas, Code Interpreter, ReAct Pattern

RAG Restaurant Ordering System

— github.com/alejortizp/rag-ordering

Conversational ordering system with RAG over restaurant menu. Agent processes product selection in natural language, generates invoices, and calculates totals. Includes inventory management.

Stack: LangChain, FAISS, Azure OpenAI, FastAPI, Streamlit

